

Biashara Bridges — Project Architecture

This document captures a professional-level architecture analysis for the Django project in this repository. It describes the current structure, responsibilities, strengths, weaknesses, identified gaps to reach "world-class" maturity, and prioritized remediation recommendations with suggested next steps.

****Location**:** `docs/ARCHITECTURE.md`

1. Executive Summary

This repository is a Django-based monolith that serves the Biashara Bridges website (landing page + additional apps). It includes multiple Django apps (`accounts`, `marketplace`, `payments`), a `config` package for Django settings, templates and static assets, service modules for payment integrations, a sqlite DB (for development), and helper scripts.

Overall maturity: medium. The codebase is functional and pragmatic: clear apps, services grouped under `payments/services`, and templates for the single-page landing site. To reach professional / world-class standards, the project requires improvements across configuration management, packaging, deployment, observability, testing, security, and modular design.

2. High-Level Architecture

- Platform: Django (WSGI/ASGI capable, `config` package contains `settings.py`, `urls.py`, `wsgi.py`, `asgi.py`).
- Apps (domain partitions):
 - `accounts` — authentication, user profiles, custom backends and forms
 - `marketplace` — marketplace models, views, templates
 - `payments` — payment models, signals, webhooks, services (stripe/paypal/mpesa), `services/` holds adapters
- Templates: server-side Django templates (`templates/` root) and numerous Element-orientated HTML artifacts (landing page). Landing page heavily references static assets.
- Static: `static/` contains CSS/JS/fonts; `static/landing/` holds downloaded/replicated assets for the landing page.
- DB: SQLite present in repo (`db.sqlite3`) — used for rapid dev/test only.
- Scripts: `scripts/\*.py` utility scripts for generating data and fetching assets.
- Integrations: third-party payment providers (Stripe, PayPal, Mpesa), webhooks, email (likely), external asset references.

Design pattern highlights:

- Services/adapters pattern inside `payments/services` is a good separation of concerns for payment integrations.
- Signals used in `payments/signals.py` for side effects.

3. File & Package Map (important files only)

- `manage.py` — Django management entry.
- `config/` — Django config (`settings.py`, `urls.py`, `wsgi.py`, `asgi.py`).
- `accounts/` — app for users/authorization.
- `marketplace/` — business logic models/views/templates.
- `payments/` — payments app, webhooks, services, signals.
- `templates/` — many site templates incl. `landing.html`.
- `static/` — local static assets.
- `db.sqlite3` — local DB (should not be used in production repo).
- `requirements.txt` — dependencies (pinning needs verification).
- `scripts/` — small utilities used for local asset fetching and data creation.

4. Strengths (what is already good)

- Logical app separation: payments, accounts, marketplace are separated into Django apps.
- Services folder for payment adapters demonstrates some clean architecture thinking (adapter/strategy pattern).
- Use of Django migrations under each app (migrations present). This is good for schema management.
- Templates & static separation follow Django conventions.
- Presence of tests files in apps (`tests.py`) — shows intent to test (coverage level TBD).
- Helper scripts exist for repeatable local tasks (data creation, asset fetching).

5. Key Weaknesses, Pain Points & Gaps (analysis)

Below are concrete pain points preventing this project from being "world-class".

1. Configuration & Secrets Management

- `config/settings.py` likely centralizes environment and secrets. No clear separation of `development` / `staging` / `production` settings or use of environment variables. Secrets may be stored in source or implicitly present via default settings.
- No documented environment file pattern (e.g., `django-environ`, `.env`) nor a secrets manager integration.

2. Deployment / Packaging

- No `Dockerfile` or containerization artifacts present (no `docker-compose.yml`).
- No CI/CD pipeline or automation for tests, linting, build, or deployment (no `.github/workflows` present).
- `db.sqlite3` checked into repository — not appropriate for production and risks leaking data.

3. Static/media management

- Large number of static assets are stored in the repo (`static/landing`) — fine for local dev but not ideal for production. No `collectstatic` strategy for S3/Cloud storage or CDN usage.
- No asset build pipeline (no webpack, esbuild, or similar) and no clear cache-busting strategy.

4. Observability & Error Monitoring

- No centralized structured logging, no Sentry/exception monitoring integration, and no healthchecks.

5. Asynchronous processing & scalability

- Payments and webhook handling are synchronous. No background job queue for long-running tasks (e.g., Celery with Redis) — webhooks and payment reconciliation may block requests.

6. Security

- No automated checks for dependencies (vulnerability scanning) or dependency pinning policies validated.
- No security hardening guidance (e.g., SSL, HSTS, secure cookies, CSP) visible in settings.

7. Testing & CI

- Tests exist but scope and coverage are unknown. No CI automation to run tests, linters, or style checks on PRs.

8. Code Hygiene & Quality

- Likely missing type annotations (mypy), linter configs (`flake8`/`ruff`) and pre-commit hooks.

- Potential duplication between services and views; no clear domain layer separation (business/domain logic probably mixed into views/models).

9. Documentation & Onboarding

- No central architecture or developer onboarding docs (this is being added). No `README` with quick-start, no `CONTRIBUTING.md`, no dev environment instructions (venv, dependencies, how to run local server securely).

10. Data & Migrations Practice

- Single `db.sqlite3` in repo may defeat testing/reproducibility and can hide migration drift. Not necessarily a problem for small projects, but for production it's a risk.

11. API & Contract Stability

- For external integrations (payments), there is likely no versioned public API or clear contract tests for webhook payloads.

6. Recommendations (prioritized)

Below are recommended actions, ordered by priority for reducing risk and moving toward world-class standards.

Top priority (immediate — 1–2 weeks):

1. Secrets & Settings Hardening

- Extract all secrets and environment-specific settings from `config/settings.py` into environment variables using `django-environ` or equivalent.
- Create `config/settings/` package with `base.py`, `development.py`, `production.py`, `local.py`. Each file imports from `base.py` and overrides env-specific values.
- Add a `.env.example` (no secrets) and document required env vars in `README`.

2. Remove `db.sqlite3` from repo & Add dev-only option

- Remove `db.sqlite3` from source control (add to `.gitignore`). If you need a dev DB, provide a script to create a throwaway sqlite or use `docker-compose` to bring up a Postgres dev DB.

3. Implement a minimal CI pipeline

- Add `.github/workflows/ci.yml` to run tests, flake/ruff linting, and safety checks on PRs. This prevents regressions and ensures basic hygiene.

4. Add dependency pinning + vulnerability checks

- Switch to pinned `requirements.txt` or use `pip-tools` (`requirements.in` -> `requirements.txt`) or `poetry` for reproducible installs.
- Add `safety` or GitHub Dependabot for automated dependency scanning.

Medium priority (2–6 weeks):

5. Containerize & provide reproducible local environment

- Add `Dockerfile` (python slim + gunicorn) and `docker-compose.yml` for local dev (Postgres, Redis for Celery). This standardizes dev and deployment.

6. Introduce background processing for payments

- Add Celery + Redis for asynchronous processing (webhook processing, reconciliation tasks). Convert blocking tasks in webhooks/signals to background jobs.

7. Static & Media strategy

- For production, use `django-storages` with S3/Cloud storage and serve static assets via CDN. Use `collectstatic` in CI/CD.
- Add an asset build step (optional) for JS/CSS bundling and optimization.

8. Observability & Error Tracking

- Integrate Sentry (or equivalent) for exception monitoring and performance traces.

- Standardize logging (structured JSON output) and expose a `/health` endpoint for monitoring.

Longer-term / formalization (6–12 weeks):

9. Improved testing & contract tests

- Expand unit and integration tests, add contract tests for payment webhooks.
- Add end-to-end smoke tests for key flows (e.g., sign-up, payments). Use GitHub Actions matrix to run tests across Python versions.

10. Code quality & developer experience

- Add pre-commit hooks (black/ruff/isort), type checking (mypy) and static code analysis.
- Add `CONTRIBUTING.md`, `README.md` with dev setup and local run instructions.

11. Architecture hardening

- Consider refactoring toward a modular/domain-driven layout: separate `services/` or `domain/` layer from web layer. Introduce well-documented interfaces between adapters (payments) and domain logic.
- For high-scale needs, evaluate breaking off critical components (payments/reconciliation, background processing) into separate services.

7. Concrete Implementation Steps (recommended quick wins)

1. Create `docs/` (done) and add this `ARCHITECTURE.md`.
2. Add `.env.example`, document `DJANGO_SECRET_KEY`, DB vars, `SENTRY_DSN`, `REDIS_URL`, `STRIPE_*`, `PAYPAL_*`, `MPESA_*`. Use `django-environ`.
3. Remove `db.sqlite3` and add `scripts/setup_dev_db.py` that can create a clear dev DB and seed minimal data.
4. Add GitHub Actions `ci.yml` to run: `pip install -r requirements.txt`, `python -m pytest`, `ruff check`, `mypy` (optional).
5. Add `Dockerfile` + `docker-compose.yml` with Postgres & Redis to reproduce environment locally.
6. Convert long-running webhook processing to Celery tasks; keep webhook endpoint fast and ack quickly.
7. Add Sentry and structured logging.
8. Add `requirements.in` + `pip-compile` to pin transitive dependencies, and enable Dependabot.

8. Security checklist (must-haves before production)

- Do NOT commit secrets. Replace any secrets in repo with env-based config.
- Set `DEBUG=False` in production and enable `ALLOWED_HOSTS` correctly.
- Use HTTPS, set `SECURE_SSL_REDIRECT`, `SESSION_COOKIE_SECURE`, `CSRF_COOKIE_SECURE`, `SECURE_HSTS_SECONDS`.
- Use a robust password hashing algorithm and enforce password policies.
- Rate-limit public endpoints (especially login & webhook endpoints) and validate webhook signatures.
- Use dependency scanning and periodic license checks.

9. Observability & SLOs

- Add Sentry for error reporting, and a Prometheus + Grafana stack or cloud alternative for metrics.
- Add basic SLOs: 99.5% uptime for site availability, 99% success for payments API calls, 95th percentile latency thresholds for page and API responses.

10. Example Roadmap (quarterly)

Quarter 1 (foundations): env & secrets, CI, remove sqlite, README, add

``.env.example`.`

Quarter 2 (ops & reliability): Docker, static/media strategy, Sentry, logging, healthchecks.

Quarter 3 (scalability): Celery & background jobs, Postgres optimisation, caching (Redis), add tests & contract tests.

Quarter 4 (polish): Type hints, stricter linting, performance optimizations, optional microservice boundary extraction.

11. Estimated Effort (very rough)

- Quick wins (CI, ``.env``, README, remove sqlite): 2–5 days
- Containerization + Docker Compose + S3 static setup: 5–10 days
- Background worker (Celery) + webhook refactor: 5–10 days
- Tests & CI coverage improvements: 2–6 weeks depending on existing test coverage

12. Offer to help

I can:

- Implement immediate quick wins: create ``Dockerfile``, ``docker-compose.yml``, ``ci.yml``, or refactor settings into ``config/settings/*.py`` and add ``.env.example``.
- Implement Celery + Redis scaffolding and convert one webhook to a background task.
- Add GitHub Actions to run tests and linting.

Tell me which priority item you want me to implement first and I will put together a plan and start implementing it.

(End of architecture document)

Immediate High-Impact Fix (Do this first)

The single fastest, highest-impact item that moves the project toward professional/world-class readiness is:

- Extracting all configuration and secrets out of ``config/settings.py`` into environment variables and reorganizing settings into environment-specific modules (development / production / local).

Why this is highest-impact:

- Prevents accidental secret/key leakage into source control.
- Enables safe CI/CD and multi-environment deployment (staging, production).
- Is a prerequisite for secure production settings (DEBUG=False, ALLOWED_HOSTS, secure cookies) and for containerization.

Clear step-by-step plan to resolve this (comprehensive, maximal)

This expanded plan is written as a practical implementation playbook. Each major step includes sub-steps, commands, verification checks, CI examples, rollbacks, and estimated time.

Major deliverables

- ``config/settings/`` package with ``base.py``, ``development.py``, ``production.py``, ``local.py``.
- ``.env.example`` and developer docs.
- ``.gitignore`` updated to exclude runtime secrets and local DB.
- A GitHub Actions workflow ``ci/secret-check.yml`` + ``ci/tests.yml`` to verify no secrets and to run tests/lint.
- A documented secret-rotation plan and optional ``scripts/rotate_secrets.md`` with steps.

Step 1 — Create a settings package (estimated 2–4 hours)

- 1.1 Create folder `config/settings/` and move `config/settings.py` -> `config/settings/base.py`.
- 1.2 Create `development.py`, `production.py`, `local.py` that each `from .base import *` then override env-specific bits.
- 1.3 Update `config/__init__.py` or the `manage.py` / `wsgi.py` / `asgi.py` environment to point `DJANGO_SETTINGS_MODULE` at a selector (e.g. `config.settings.development` for local). Use an environment variable `DJANGO_ENV` to choose automatically in deployment scripts.

Example layout:

```

...
config/
settings/
__init__.py # loads settings according to DJANGO_ENV or fallback
base.py
development.py
production.py
local.py
...

```

Snippet for `config/settings/__init__.py`:

```

```python
import os
env = os.getenv("DJANGO_ENV", "development")
if env == "production":
 from .production import *
elif env == "local":
 from .local import *
else:
 from .development import *
```

```

Why and verification

- Why: This enables per-environment configuration and makes production-only settings clearly visible.
- Verify: `python manage.py check --deploy` (after setting env) and ensure no sensitive defaults remain.

Step 2 — Add environment loader & define required variables (estimated 1–2 hours)

- 2.1 Add `django-environ` to `requirements.txt` (or `python-dotenv` if preferred).
- 2.2 In `base.py` load environment at the top and declare all required variables and defaults. Use `env.bool()` for flags.

Example snippet for `base.py` using `django-environ`:

```

```python
import environ
env = environ.Env(
 DEBUG=(bool, False),
)
environ.Env.read_env() # reads .env for local

SECRET_KEY = env("DJANGO_SECRET_KEY")
DEBUG = env.bool("DEBUG", default=False)
DATABASES = {"default": env.db()}
ALLOWED_HOSTS = env.list("ALLOWED_HOSTS", default=["localhost"])
```

```

- 2.3 Create a declarative list `REQUIRED_ENV_VARS` and at startup assert they are present in production mode. Example check inside `base.py`:

```

```python

```

```

if not DEBUG:
missing = [k for k in REQUIRED_ENV_VARS if not env(k, default=None)]
if missing:
raise RuntimeError("Missing required env vars: %s" % ", ".join(missing))

```

#### Why and verification

- Why: Centralized env loading enforces consistent parsing of DB URLs, booleans, lists and reduces ad-hoc reading of `os.environ` across the codebase.
- Verify: run ``python -c "import django; print('env ok')"`` after setting minimal env vars.

#### Step 3 — Provide ``.env.example`` and update docs (estimated 30–60 minutes)

3.1 Add ``.env.example`` at repo root with annotated keys (no secrets). Suggested keys:

```

...
Django
DJANGO_ENV=development
DJANGO_SECRET_KEY=
DEBUG=True
ALLOWED_HOSTS=127.0.0.1,localhost

Database
DATABASE_URL=sqlite:///./db.sqlite3

Sentry / Observability
SENTRY_DSN=

Redis / Celery
REDIS_URL=redis://localhost:6379/0

Payment providers
STRIPE_SECRET=
PAYPAL_CLIENT_ID=
PAYPAL_SECRET=
MPESA_KEY=

```

3.2 Document in `README.md`` the exact commands to create ``.env`` from ``.env.example`` and how to run the dev server.

#### Why and verification

- Why: Ensures new contributors can bootstrap quickly and consistently.
- Verify: Create ``.env`` locally from ``.env.example`` and run ``python manage.py runserver``.

#### Step 4 — Prevent secrets from being committed & clean history if required (estimated 1–3 hours)

4.1 Update ``.gitignore`` to include:

```

...
.env
db.sqlite3
*.py[cod]
__pycache__/_

```

4.2 If secrets already exist in git history:

- Identify sensitive commits with ``git log -S "SECRET_KEY"`` or use ``git-secrets`` to scan.
- Rotate any secrets exposed (API keys, DB passwords) immediately with their providers.
- If necessary, rewrite history with ``git filter-repo`` or ``bfg`` and coordinate the team to re-clone. Document the steps and perform during a maintenance window.

## Why and verification

- Why: Adds a safety net preventing future accidental commits and reduces the impact window of leaked credentials.
- Verify: ``git ls-files --others --exclude-standard`` should not show ``.env`` or ``db.sqlite3``.

## Step 5 — Add CI checks to catch regressions (estimated 2–6 hours)

### 5.1 Add GitHub Actions workflow ``ci/secret-check.yml`` that:

- Runs a tiny Python script or ``grep`` to assert that production settings file(s) contain no literal ``SECRET_KEY`` assignments or hard-coded tokens.
- Runs ``pytest -q`` and ``ruff`` lint.

Example ``ci/secret-check.yml`` skeleton:

```
``yaml
name: CI
on: [push, pull_request]
jobs:
 checks:
 runs-on: ubuntu-latest
 steps:
 - uses: actions/checkout@v4
 - name: Set up Python
 uses: actions/setup-python@v4
 with:
 python-version: '3.11'
 - name: Install deps
 run: |
 python -m pip install --upgrade pip
 pip install -r requirements.txt
 - name: Secret scan (simple)
 run: |
 if grep -R "SECRET_KEY\s*=\s*['\"]" config | grep -v "__init__"; then echo
 "Found SECRET_KEY in repo"; exit 1; fi
 - name: Run tests
 run: pytest -q
 - name: Lint
 run: ruff check .
````
```

5.2 Optionally add ``git-secrets`` or use a GitHub Action for secret scanning (e.g., ``gitleaks-action``).

Why and verification

- Why: CI gates minimize the chance of secret or config regressions reaching main and ensure tests/linting run on every PR.
- Verify: Create a test commit that intentionally includes a ``SECRET_KEY`` in a non-excluded file and ensure CI fails.

Step 6 — Verification checklist & acceptance criteria (automated + manual) (estimated 1–2 hours)

- ``config/settings/production.py`` contains zero hard-coded credentials.
- ``.env.example`` present at repo root and ``.env`` does not exist in repo.
- ``.gitignore`` excludes ``.env`` and ``db.sqlite3`` and no secrets appear in changed files in PRs.
- CI workflow runs on PRs and fails if a ``SECRET_KEY`` literal is present.
- Dev run: contributor can follow README, create ``.env`` from ``.env.example``, and start ``manage.py runserver`` without secret errors.

Automated checks to add (recommended):

- ``pytest`` for unit tests.
- ``ruff`` for linting.
- ``gitleaks`` or ``gitleaks-action`` for secret scanning.

Step 7 — Secret rotation & incident response (must-have if secrets were committed)

7.1 Immediately rotate all compromised keys in respective provider consoles (Stripe, PayPal, hosting, Sentry).

7.2 Revoke old keys and issue new ones, update `.env` for deployments.

7.3 If a secret was in git history, run `git filter-repo --path-glob "*" --replace-refs delete-no-add` or BFG as documented; coordinate with the team to re-clone and reset CI tokens.

7.4 Document the event in `SECURITY.md` and notify stakeholders.

Step 8 — Rollout & rollback plan (minimal disruption)

8.1 Implement changes in a feature branch; open a PR and request at least one reviewer.

8.2 Merge after CI passes and after manual verification on a staging environment.

8.3 Rollback: if production shows issues, revert the deployment and reopen the PR, or restore old environment variables from secret manager.

Step 9 — Timeline & responsibility

- Phase A (1–3 days): Implement settings package, `.env.example`, update `.gitignore`, add minimal README snippets.

- Phase B (2–5 days): Add CI workflows, linting, run tests, add simple secret scanning action.

- Phase C (2–7 days): If required, rotate leaked secrets and clean git history; integrate secret manager (HashiCorp Vault / AWS Secrets Manager) for production.

Step 10 — Follow-ups (post completion)

- Containerize (Docker + `docker-compose`) and switch development to Postgres + Redis for realistic parity with production.

- Add Sentry and structured logging for observability.

- Convert webhook handling to background tasks (Celery + Redis) and add contract tests for payment webhooks.

Risk & mitigation summary

- Risk: Rewriting git history when removing secrets is destructive. Mitigation: perform during a scheduled maintenance window, rotate secrets, and communicate the change.

- Risk: CI false positives. Mitigation: start with conservative checks (grep for `SECRET_KEY`) and iterate.

Acceptance criteria (final)

- No secrets in `config` files or committed history (or otherwise rotated and removed from history).

- Clear developer onboarding: `.env.example`, README updated, and quick start commands working.

- CI enforces the rules on PRs and merges.

Ready-to-run commands (PowerShell)

```
```powershell
create local env
copy .env.example .env
install deps
venv\Scripts\Activate; pip install -r requirements.txt
run locally
python manage.py migrate
python manage.py runserver 127.0.0.1:8000
```
```

This completes the comprehensive immediate-fix playbook. Execute the steps in small commits and validate each automated check before moving to the next.

(End of added immediate-fix plan)